

Mike Walker - Term Project

```
In [1]: import seaborn as sns
import numpy as np
import PIL
import pathlib

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras import regularizers
import tensorflow.keras.layers.experimental.preprocessing as tfe
from tensorflow.keras.models import Sequential

print(tf.__version__)
```

2.6.0

```
In [2]: def show_picture(noun):
    preview_image = list(data_dir.glob(noun + '/*'))
    preview_image_url = str(preview_image[0])
    display(PIL.Image.open(preview_image_url))
```

Activate GPU

```
In [3]: #Activate and select the GPU
my_desktop_gpus = tf.config.experimental.list_physical_devices('GPU')
if len(my_desktop_gpus) > 0:
    for gpu in my_desktop_gpus:
        tf.config.experimental.set_memory_growth(gpu, False)
        print(f'{gpu} memory growth: {tf.config.experimental.get_memory_growth(gpu)}
```

PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU') memory growth: False

Select Photo Set

```
In [4]: test_examples = {
    "family_padded": r"padded_family_photos",
    "family_no_pad": r"family_photos",
    "family_tiny_pad": r"padded_portrait_family",
    "flowers": r"flower_photos"
}

folder_url = test_examples["flowers"]
print(folder_url)


data_dir = pathlib.Path(folder_url)

image_count = len(list(data_dir.glob('*/*')))
print(image_count)
```

flower_photos
3665

Test Photo Set Loaded Correctly

```
In [5]: if (folder_url == test_examples["flowers"]):  
        show_picture("daisy")  
        show_picture("dandelion")  
        show_picture("roses")  
        show_picture("sunflowers")  
        show_picture("tulips")  
    else:  
        show_picture("london")  
        show_picture("mike")  
        show_picture("robin")
```





In [6]: `### Set image and model parameters`

```
In [7]: if (folder_url == test_examples["family_tiny_pad"]):  
        batch_size = 4 #had to adjust for large images  
        img_height = 2049  
        img_width = 1539  
    elif (folder_url == test_examples["flowers"]):  
        batch_size = 64
```

```

img_height = 240
img_width = 180
elif (folder_url == test_examples["family_no_pad"]):
    batch_size = 4
    img_height = 2048 #not actual size, just reduced it
    img_width = 1536 #not actual size, just reduced it
elif (folder_url == test_examples["family_padded"]):
    #This example has a massive padding.
    batch_size = 4 #had to adjust for large images
    img_height = 2305
    img_width = 2100

N_TRAIN = int(1e4)
STEPS_PER_EPOCH = N_TRAIN//batch_size

num_of_epochs = [60]

```

Load Training Dataset and Validation Dataset

```

In [8]: train_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="training",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)

```

Found 3665 files belonging to 5 classes.
Using 2932 files for training.

```

In [9]: val_ds = tf.keras.utils.image_dataset_from_directory(
    data_dir,
    validation_split=0.2,
    subset="validation",
    seed=123,
    image_size=(img_height, img_width),
    batch_size=batch_size)

```

Found 3665 files belonging to 5 classes.
Using 733 files for validation.

Verify Output and Shape

```

In [10]: class_names = train_ds.class_names
    print(class_names)

```

```
['daisy', 'dandelion', 'roses', 'sunflowers', 'tulips']
```

```

In [11]: for image_batch, labels_batch in train_ds:
    print(image_batch.shape)
    print(labels_batch.shape)
    break

```

```
(64, 240, 180, 3)
(64,)
```

Setup and Run Model

```
In [12]: AUTOTUNE = tf.data.AUTOTUNE
```

```
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=AUTOTUNE)
val_ds = val_ds.cache().prefetch(buffer_size=AUTOTUNE)
```

```
In [13]: normalization_layer = tfe.Rescaling(1./255)
```

```
In [14]: normalized_ds = train_ds.map(lambda x, y: (normalization_layer(x), y))
image_batch, labels_batch = next(iter(normalized_ds))
first_image = image_batch[0]
# Notice the pixel values are now in `[0,1]`.
print(np.min(first_image), np.max(first_image))
```

```
0.0 1.0
```

```
In [15]: # Improving the Learning model by introducing variations in the training data
data_augmentation = Sequential([
    tfe.RandomFlip("horizontal", input_shape=(img_height, img_width, 3)),
    tfe.RandomRotation(0.1),
    tfe.RandomZoom(0.1),
])
```

```
In [16]: #model inspried by: https://www.tensorflow.org/tutorials/keras/overfit_and_underfit
num_classes = len(class_names) #outputs
```

```
model = Sequential([
    data_augmentation,
    tfe.Rescaling(1./255),
    layers.Conv2D(16, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(32, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Conv2D(64, 3, padding='same', activation='relu'),
    layers.MaxPooling2D(),
    layers.Dropout(0.2),
    layers.Flatten(),
    layers.Dense(128, activation='relu'),
    layers.Dense(num_classes, name="outputs")
])
```

```
#epochs are high (50 and in one case 500) so adding this to help improve the conver
#by gradually reducing the learning rate as training progresses (as stated in Tenso
#from what I understood, it help led to a more stable/faster convergence to a good
lr_schedule = tf.keras.optimizers.schedules.InverseTimeDecay(
    0.001,
    decay_steps=STEPS_PER_EPOCH*1000,
    decay_rate=1,
    staircase=False)
```

```
#Selected Adam for the optimizer
```

```
optimizer = tf.keras.optimizers.Adam(learning_rate=lr_schedule)

model.compile(optimizer=optimizer,
              loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
              metrics=['accuracy'])

model.summary()

history = model.fit(train_ds, validation_data = val_ds, epochs = num_of_epochs[0])

results_data = {'Training Accuracy': history.history['accuracy'], 'Validation Accuracy': history.history['val_accuracy']}

history = None
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
sequential (Sequential)	(None, 240, 180, 3)	0

rescaling_1 (Rescaling)	(None, 240, 180, 3)	0

conv2d (Conv2D)	(None, 240, 180, 16)	448

max_pooling2d (MaxPooling2D)	(None, 120, 90, 16)	0

conv2d_1 (Conv2D)	(None, 120, 90, 32)	4640

max_pooling2d_1 (MaxPooling2D)	(None, 60, 45, 32)	0

conv2d_2 (Conv2D)	(None, 60, 45, 64)	18496

max_pooling2d_2 (MaxPooling2D)	(None, 30, 22, 64)	0

dropout (Dropout)	(None, 30, 22, 64)	0

flatten (Flatten)	(None, 42240)	0

dense (Dense)	(None, 128)	5406848

outputs (Dense)	(None, 5)	645
=====		
Total params: 5,431,077		
Trainable params: 5,431,077		
Non-trainable params: 0		

Epoch 1/60

46/46 [=====] - 5s 33ms/step - loss: 1.6113 - accuracy: 0.2797 - val_loss: 1.2949 - val_accuracy: 0.4502

Epoch 2/60

46/46 [=====] - 1s 17ms/step - loss: 1.2521 - accuracy: 0.4894 - val_loss: 1.0494 - val_accuracy: 0.5771

Epoch 3/60

46/46 [=====] - 1s 17ms/step - loss: 1.0589 - accuracy: 0.5754 - val_loss: 1.0028 - val_accuracy: 0.6221

Epoch 4/60

46/46 [=====] - 1s 17ms/step - loss: 0.9670 - accuracy: 0.6211 - val_loss: 0.9947 - val_accuracy: 0.6221

Epoch 5/60

46/46 [=====] - 1s 17ms/step - loss: 0.9057 - accuracy: 0.6429 - val_loss: 0.8860 - val_accuracy: 0.6658

Epoch 6/60

46/46 [=====] - 1s 17ms/step - loss: 0.8618 - accuracy: 0.6675 - val_loss: 0.9322 - val_accuracy: 0.6576

Epoch 7/60

46/46 [=====] - 1s 17ms/step - loss: 0.8143 - accuracy: 0.6695 - val_loss: 0.8531 - val_accuracy: 0.6985

Epoch 8/60

46/46 [=====] - 1s 17ms/step - loss: 0.7850 - accuracy: 0.6883 - val_loss: 0.8606 - val_accuracy: 0.6767

Epoch 9/60
46/46 [=====] - 1s 17ms/step - loss: 0.7729 - accuracy: 0.7074 - val_loss: 0.8477 - val_accuracy: 0.6671
Epoch 10/60
46/46 [=====] - 1s 17ms/step - loss: 0.7278 - accuracy: 0.7173 - val_loss: 0.7805 - val_accuracy: 0.7162
Epoch 11/60
46/46 [=====] - 1s 17ms/step - loss: 0.6880 - accuracy: 0.7306 - val_loss: 0.7422 - val_accuracy: 0.7312
Epoch 12/60
46/46 [=====] - 1s 17ms/step - loss: 0.6501 - accuracy: 0.7514 - val_loss: 0.8251 - val_accuracy: 0.7135
Epoch 13/60
46/46 [=====] - 1s 17ms/step - loss: 0.6466 - accuracy: 0.7538 - val_loss: 0.7954 - val_accuracy: 0.7190
Epoch 14/60
46/46 [=====] - 1s 17ms/step - loss: 0.6329 - accuracy: 0.7568 - val_loss: 0.7590 - val_accuracy: 0.7312
Epoch 15/60
46/46 [=====] - 1s 17ms/step - loss: 0.6127 - accuracy: 0.7650 - val_loss: 0.8522 - val_accuracy: 0.6944
Epoch 16/60
46/46 [=====] - 1s 17ms/step - loss: 0.6106 - accuracy: 0.7742 - val_loss: 0.7598 - val_accuracy: 0.7435
Epoch 17/60
46/46 [=====] - 1s 17ms/step - loss: 0.5536 - accuracy: 0.7841 - val_loss: 0.7441 - val_accuracy: 0.7422
Epoch 18/60
46/46 [=====] - 1s 17ms/step - loss: 0.5514 - accuracy: 0.7902 - val_loss: 0.7596 - val_accuracy: 0.7190
Epoch 19/60
46/46 [=====] - 1s 17ms/step - loss: 0.5222 - accuracy: 0.8035 - val_loss: 0.7739 - val_accuracy: 0.7517
Epoch 20/60
46/46 [=====] - 1s 17ms/step - loss: 0.5138 - accuracy: 0.8087 - val_loss: 0.7333 - val_accuracy: 0.7326
Epoch 21/60
46/46 [=====] - 1s 17ms/step - loss: 0.5004 - accuracy: 0.8114 - val_loss: 0.7403 - val_accuracy: 0.7367
Epoch 22/60
46/46 [=====] - 1s 17ms/step - loss: 0.4789 - accuracy: 0.8186 - val_loss: 0.8192 - val_accuracy: 0.7231
Epoch 23/60
46/46 [=====] - 1s 17ms/step - loss: 0.4708 - accuracy: 0.8247 - val_loss: 0.8391 - val_accuracy: 0.7353
Epoch 24/60
46/46 [=====] - 1s 17ms/step - loss: 0.4512 - accuracy: 0.8336 - val_loss: 0.7910 - val_accuracy: 0.7285
Epoch 25/60
46/46 [=====] - 1s 17ms/step - loss: 0.4365 - accuracy: 0.8366 - val_loss: 0.7867 - val_accuracy: 0.7435
Epoch 26/60
46/46 [=====] - 1s 17ms/step - loss: 0.4181 - accuracy: 0.8496 - val_loss: 0.7752 - val_accuracy: 0.7326
Epoch 27/60
46/46 [=====] - 1s 17ms/step - loss: 0.3930 - accuracy: 0.8

554 - val_loss: 0.7926 - val_accuracy: 0.7340
Epoch 28/60
46/46 [=====] - 1s 17ms/step - loss: 0.3945 - accuracy: 0.8
591 - val_loss: 0.7723 - val_accuracy: 0.7462
Epoch 29/60
46/46 [=====] - 1s 17ms/step - loss: 0.3932 - accuracy: 0.8
540 - val_loss: 0.8463 - val_accuracy: 0.7353
Epoch 30/60
46/46 [=====] - 1s 17ms/step - loss: 0.3770 - accuracy: 0.8
550 - val_loss: 0.8277 - val_accuracy: 0.7381
Epoch 31/60
46/46 [=====] - 1s 16ms/step - loss: 0.3362 - accuracy: 0.8
806 - val_loss: 0.8109 - val_accuracy: 0.7449
Epoch 32/60
46/46 [=====] - 1s 17ms/step - loss: 0.3552 - accuracy: 0.8
704 - val_loss: 0.7971 - val_accuracy: 0.7503
Epoch 33/60
46/46 [=====] - 1s 17ms/step - loss: 0.3276 - accuracy: 0.8
840 - val_loss: 0.8420 - val_accuracy: 0.7340
Epoch 34/60
46/46 [=====] - 1s 17ms/step - loss: 0.2936 - accuracy: 0.8
932 - val_loss: 0.8728 - val_accuracy: 0.7449
Epoch 35/60
46/46 [=====] - 1s 17ms/step - loss: 0.2930 - accuracy: 0.8
898 - val_loss: 0.8690 - val_accuracy: 0.7476
Epoch 36/60
46/46 [=====] - 1s 17ms/step - loss: 0.3136 - accuracy: 0.8
844 - val_loss: 0.8333 - val_accuracy: 0.7462
Epoch 37/60
46/46 [=====] - 1s 17ms/step - loss: 0.2762 - accuracy: 0.9
059 - val_loss: 0.8426 - val_accuracy: 0.7640
Epoch 38/60
46/46 [=====] - 1s 17ms/step - loss: 0.2612 - accuracy: 0.9
048 - val_loss: 0.9501 - val_accuracy: 0.7394
Epoch 39/60
46/46 [=====] - 1s 17ms/step - loss: 0.2708 - accuracy: 0.9
052 - val_loss: 1.0670 - val_accuracy: 0.7367
Epoch 40/60
46/46 [=====] - 1s 17ms/step - loss: 0.2686 - accuracy: 0.9
072 - val_loss: 1.1628 - val_accuracy: 0.7299
Epoch 41/60
46/46 [=====] - 1s 17ms/step - loss: 0.3146 - accuracy: 0.8
885 - val_loss: 0.9346 - val_accuracy: 0.7271
Epoch 42/60
46/46 [=====] - 1s 17ms/step - loss: 0.2659 - accuracy: 0.9
025 - val_loss: 0.8501 - val_accuracy: 0.7667
Epoch 43/60
46/46 [=====] - 1s 17ms/step - loss: 0.2212 - accuracy: 0.9
192 - val_loss: 0.9953 - val_accuracy: 0.7353
Epoch 44/60
46/46 [=====] - 1s 17ms/step - loss: 0.2226 - accuracy: 0.9
154 - val_loss: 0.9358 - val_accuracy: 0.7667
Epoch 45/60
46/46 [=====] - 1s 17ms/step - loss: 0.2208 - accuracy: 0.9
188 - val_loss: 0.9009 - val_accuracy: 0.7544
Epoch 46/60


```

46/46 [=====] - 1s 17ms/step - loss: 0.2169 - accuracy: 0.9
253 - val_loss: 0.9819 - val_accuracy: 0.7490
Epoch 47/60
46/46 [=====] - 1s 17ms/step - loss: 0.1910 - accuracy: 0.9
294 - val_loss: 0.9992 - val_accuracy: 0.7572
Epoch 48/60
46/46 [=====] - 1s 17ms/step - loss: 0.2285 - accuracy: 0.9
233 - val_loss: 1.0239 - val_accuracy: 0.7531
Epoch 49/60
46/46 [=====] - 1s 17ms/step - loss: 0.1873 - accuracy: 0.9
345 - val_loss: 1.1042 - val_accuracy: 0.7572
Epoch 50/60
46/46 [=====] - 1s 17ms/step - loss: 0.2028 - accuracy: 0.9
321 - val_loss: 0.9734 - val_accuracy: 0.7585
Epoch 51/60
46/46 [=====] - 1s 17ms/step - loss: 0.1668 - accuracy: 0.9
393 - val_loss: 1.1072 - val_accuracy: 0.7531
Epoch 52/60
46/46 [=====] - 1s 17ms/step - loss: 0.1936 - accuracy: 0.9
294 - val_loss: 1.0176 - val_accuracy: 0.7599
Epoch 53/60
46/46 [=====] - 1s 17ms/step - loss: 0.1500 - accuracy: 0.9
495 - val_loss: 1.0397 - val_accuracy: 0.7599
Epoch 54/60
46/46 [=====] - 1s 17ms/step - loss: 0.1534 - accuracy: 0.9
478 - val_loss: 1.1137 - val_accuracy: 0.7490
Epoch 55/60
46/46 [=====] - 1s 17ms/step - loss: 0.1472 - accuracy: 0.9
492 - val_loss: 1.1271 - val_accuracy: 0.7572
Epoch 56/60
46/46 [=====] - 1s 17ms/step - loss: 0.1416 - accuracy: 0.9
492 - val_loss: 1.2131 - val_accuracy: 0.7435
Epoch 57/60
46/46 [=====] - 1s 17ms/step - loss: 0.1529 - accuracy: 0.9
482 - val_loss: 1.2705 - val_accuracy: 0.7326
Epoch 58/60
46/46 [=====] - 1s 17ms/step - loss: 0.1631 - accuracy: 0.9
389 - val_loss: 1.2168 - val_accuracy: 0.7572
Epoch 59/60
46/46 [=====] - 1s 17ms/step - loss: 0.1353 - accuracy: 0.9
505 - val_loss: 1.2320 - val_accuracy: 0.7544
Epoch 60/60
46/46 [=====] - 1s 17ms/step - loss: 0.1257 - accuracy: 0.9
557 - val_loss: 1.2702 - val_accuracy: 0.7490

```

```
In [17]: sns.lineplot(data=results_data, dashes=False).set(title='Training vs Validation Acc
```

```
Out[17]: [Text(0.5, 1.0, 'Training vs Validation Accuracy')]
```



```
In [20]: if (folder_url == test_examples["family_tiny_pad"]):
    london = r"test_family_photos\tiny_padded_test\tiny-pad_london_test_1.jpg"
    mike = r"test_family_photos\tiny_padded_test\tiny-pad_mike_test_1.jpg"
    robin = r"test_family_photos\tiny_padded_test\tiny-pad_robin_test_1.jpg"

    test_images = {"london": london, "mike": mike, "robin": robin}

elif (folder_url == test_examples["flowers"]):
    daisy = r"test_flower_photos\daisy_test.jpg"
    dandelion = r"test_flower_photos\dandelion_test.jpg"
    rose = r"test_flower_photos\rose_test.jpg"
    sunflower = r"test_flower_photos\sunflower_test.jpg"
    tulip = r"test_flower_photos\tulip_test.jpg"

    test_images = {"daisy": daisy, "dandelion": dandelion, "rose": rose, "sunflower": sunflower, "tulip": tulip}

elif (folder_url == test_examples["family_no_pad"]):
    london = r"test_family_photos\no_padding_test\no-pad_london_test_1.jpg"
    mike = r"test_family_photos\no_padding_test\no-pad_mike_test_1.jpg"
    robin = r"test_family_photos\no_padding_test\no-pad_robin_test_1.jpg"

    test_images = {"london": london, "mike": mike, "robin": robin}

elif (folder_url == test_examples["family_padded"]):
    london = r"test_family_photos\large_padded_test\padded_london_test_1.jpg"
    mike = r"test_family_photos\large_padded_test\padded_mike_test_1.jpg"
    robin = r"test_family_photos\large_padded_test\padded_robin_test_1.jpg"

    test_images = {"london": london, "mike": mike, "robin": robin}
```

```

for test_image in test_images:
    print("Testing with " + test_image)

    final_test_img = tf.keras.utils.load_img(test_images[test_image], target_size=(
img_array = tf.keras.utils.img_to_array(final_test_img)
img_array = tf.expand_dims(img_array, 0) # Create a batch

    predictions = model.predict(img_array)
    score = tf.nn.relu(predictions[0])

    print(
        "This image most likely belongs to {} with a {:.2f} percent confidence. \n"
        .format(class_names[np.argmax(score)], 100 * np.max(score))
    )

    img = None
    img_array = None
    predictions = None
    score = None

```

Testing with daisy

This image most likely belongs to daisy with a 1383.15 percent confidence.

Testing with dandelion

This image most likely belongs to dandelion with a 649.88 percent confidence.

Testing with rose

This image most likely belongs to roses with a 1333.58 percent confidence.

Testing with sunflower

This image most likely belongs to sunflowers with a 793.11 percent confidence.

Testing with tulip

This image most likely belongs to tulips with a 1468.01 percent confidence.

In [19]:

```

for test_image in test_images:
    print("test photo: " + test_image)
    display(PIL.Image.open(test_images[test_image]))
    print()

```

test photo: daisy



test photo: dandelion



test photo: rose



test photo: sunflower



test photo: tulip

